


## Workflow 5 PDF

Initial Network:

```
super().__init__(**kwargs)
Model: "sequential"
```

Layer (type)	Output Shape	Param #
rescaling (Rescaling)	(None, 150, 150, 3)	0
max_pooling2d (MaxPooling2D)	(None, 75, 75, 3)	0
conv2d (Conv2D)	(None, 75, 75, 32)	896
max_pooling2d_1 (MaxPooling2D)	(None, 37, 37, 32)	0
conv2d_1 (Conv2D)	(None, 37, 37, 48)	13,872
max_pooling2d_2 (MaxPooling2D)	(None, 18, 18, 48)	0
conv2d_2 (Conv2D)	(None, 18, 18, 48)	20,784
max_pooling2d_3 (MaxPooling2D)	(None, 9, 9, 48)	0
flatten (Flatten)	(None, 3888)	0
dense (Dense)	(None, 25)	97,225
dense_1 (Dense)	(None, 3)	78

Total params: 132,855 (518.96 KB)  
Trainable params: 132,855 (518.96 KB)  
Non-trainable params: 0 (0.00 B)  
Epoch 1/9



The accuracy and loss of our best-learned model (which was at epoch 9 before overfitting) were 74.92% (0.7492) and 64.06% (0.6406) respectively. The final trained model at epoch 20 got an accuracy of 71.38% (0.7138) and a loss of 74.06% (0.7406), confirming that overfitting begins after epoch 9.